

## Week 1 & 2

### OpenCV Installation Command:

```
| pip install opencv-python
```

### Importing an Image

```
import cv2  
from google.colab.patches import cv2_imshow  
img = cv2.imread("Lena.png")  cv2_imshow(img)
```



### Converting image into RGB

```
import cv2  import  
matplotlib.pyplot as plt img =  
cv2.imread("Lena.png")  
cv2_imshow(img)
```



```
image_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB) cv2.imshow(img)
```



## Converting image into GRAY

```
import cv2 import
matplotlib.pyplot as plt img =
cv2.imread("Lena.png")
print("BGR")
cv2.imshow(img)
    print("GRAY")
    image_gray =
cv2.cvtColor(img,
```

```
cv2.COLOR_BGR2GRAY)  
cv2.imshow(image_gray)
```

BGR



GRAY



## Plot the three channels of the image

```
import numpy as np  
img = cv2.imread("Lena.png", cv2.IMREAD_UNCHANGED)
```

## Blue Channel

```
blue_channel = img[:, :, 0] blue_img  
= np.zeros(img.shape)  
blue_img[:, :, 0] = blue_channel  
cv2_imshow(blue_img)
```



## Green Channel

```
green_channel = img[:, :, 1] green_img  
= np.zeros(img.shape)  
green_img[:, :, 1] = green_channel  
cv2_imshow(green_img)
```



## Red Channel

```
red_channel = img[:, :, 2] red_img  
= np.zeros(img.shape)  
red_img[:, :, 2] = red_channel  
cv2_imshow(red_img)
```



## Transform the image into HLS

```
img = cv2.imread("Lena.png")  img_hls =  
cv2.cvtColor(img, cv2.COLOR_BGR2HLS)  
cv2_imshow(img_hls)
```





## Transform the image into HSV


```
img = cv2.imread("Lena.png")    img_hsv =  
cv2.cvtColor(img, cv2.COLOR_BGR2HSV)  
cv2.imshow(img_hsv)
```



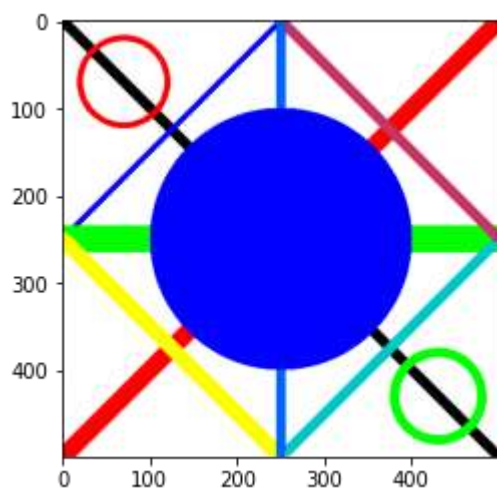
## Drawing rectangle, lines, circle on Image

```
img = np.zeros((500,500,3), dtype="uint8")  
# Changing the color of the image  
img[:] = (255,255,255) # Drawing  
a black line  
cv2.line(img, (0,0), (500,500), (0,0,0), (10))  
# Drawing a blue line  
cv2.line(img, (0,500), (500,0), (255,0,0), (20))  
# Drawing a red line  
cv2.line(img, (0,250), (250,0), (0,0,255), (5))  
# Drawing a green line  
cv2.line(img, (500,250), (0,250), (0,255,0), (30))  
# Drawing a yellow line  
cv2.line(img, (500,250), (250,500), (0,200,200), (10))  
# Drawing a violet line  
cv2.line(img, (250,0), (500,250), (200,50,100), (10))  
# Drawing a cyan line  
cv2.line(img, (0,250), (250,500), (255,255,0), (15))  
# Drawing an orange line  
cv2.line(img, (250,0), (250,500), (0,100,255), (10))
```

```
plt.imshow(img)
<matplotlib.image.AxesImage
age at 0x7ff65415c710>
```




```
cv2.circle(img, (250,250), 150, (0,0,255), (-1))
# Drawing a blue circle
cv2.circle(img, (70,70), 50, (255,0,0), (5))
# Drawing a green circle
cv2.circle(img, (430,430), 50, (0,255,0), (10))
plt.imshow(img) <matplotlib.image.AxesImage at
0x7ff654124e90>
```



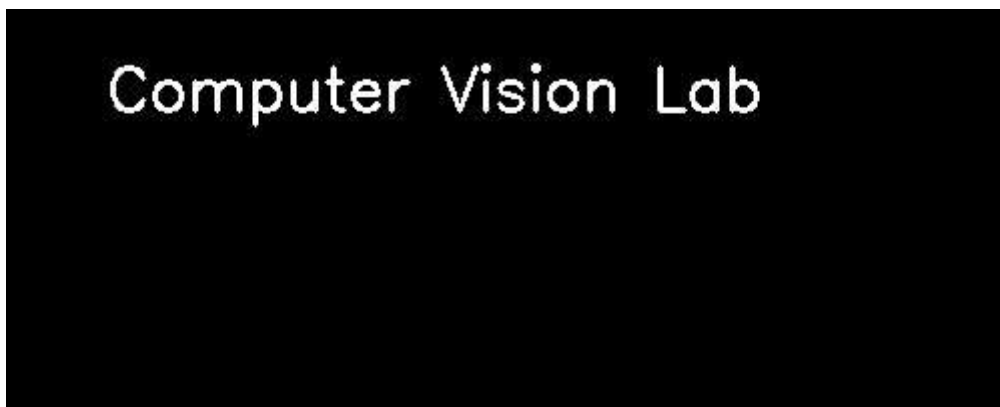
```
img = np.zeros((500,500,3),dtype='uint8')
image = cv2.rectangle(img, (25,25), (400,400), (0,255,0), 7)
plt.imshow(image)

<matplotlib.image.AxesImage at 0x7ff653f138d0>
```



## Adding Text on Image

```
image = np.zeros((200,500))
font = cv2.FONT_HERSHEY_SIMPLEX
org = (50, 50)    fontStyle =
1    color = (255, 0, 0)
thickness = 2
image = cv2.putText(image, 'Computer Vision Lab', org, font,
fontStyle, (255,0,0), thickness, cv2.LINE_AA) cv2.imshow(image)
```



## Week 3 & 4

### Storing and Accessing many Images using Python (pickle)

```
from PIL import Image    import
pickle
```




```
img = Image.open('Lena.png') with  
open('savedimage.pkl', 'wb') as f:  
pickle.dump(img, f)
```

## Week 5 & 6


# Image Segmentation Using Color Spaces in OpenCV + Python

## Color Spaces and Reading Images in OpenCV


```
img = cv2.imread("cmyk_paint.png", cv2.IMREAD_UNCHANGED)  
cv2_imshow(img)
```




```
img = cv2.imread("cmyk_paint.png", cv2.IMREAD_UNCHANGED)  
blue_channel = img[:, :, 0] blue_img = np.zeros(img.shape)  
blue_img[:, :, 0] = blue_channel print("Blue Color Space")  
cv2_imshow(blue_img) Blue color space
```



```
img = cv2.imread("cmyk_paint.png", cv2.IMREAD_UNCHANGED)
green_channel = img[:, :, 1] green_img =
np.zeros(img.shape) green_img[:, :, 1] = green_channel
print("Green Color Space")      cv2_imshow(green_img)
Green Color Space
```




```
img = cv2.imread("cmyk_paint.png", cv2.IMREAD_UNCHANGED)
red_channel = img[:, :, 2] red_img = np.zeros(img.shape)
red_img[:, :, 2] = red_channel print("Red Color Space")
cv2_imshow(red_img) Red Color Space
```




# Visualizing Nemo in RGB Color Space

```
nemo_img = cv2.imread("nemo5.jpg") cv2.imshow(nemo_img)
```



```
print("Visualizing Nemo in RGB Color Space")  
nemo_rgb = cv2.cvtColor(nemo_img, cv2.COLOR_BGR2RGB)  
cv2.imshow(nemo_rgb)
```

Visualizing Nemo in RGB Color Space




# Visualizing Nemo in HSV Color Space

```
print("Visualizing Nemo in HSV Color Space")    nemo_hsv
= cv2.cvtColor(nemo_img, cv2.COLOR_BGR2HSV)
cv2.imshow(nemo_hsv)
```

Visualizing Nemo in HSV Color Space



```
from mpl_toolkits.mplot3d import Axes3D
from matplotlib import cm    from
matplotlib import colors
    r, g, b =
cv2.split(nemo_img) fig =
plt.figure()
axis = fig.add_subplot(1, 1, 1, projection="3d")
pixel_colors = nemo_img.reshape((np.shape(nemo_img)
    [0]*np.shape(nemo_img)[1],
        3)) norm =
colors.Normalize(vmin=-1.,vmax=1.)
norm.autoscale(pixel_colors) pixel_colors =
norm(pixel_colors).tolist()
axis.scatter(r.flatten(), g.flatten(), b.flatten(),
    facecolors=pixel_colors, marker=".")
axis.set_xlabel("Red") axis.set_ylabel("Green")
axis.set_zlabel("Blue") plt.show()
```




## Week 7 & 8

### Face Detection in Python, OpenCV

```
face_cascade =
    cv2.CascadeClassifier('haarcascade_frontalface_default.xml')
img = cv2.imread("Lena.png")    gray =
cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)    faces =
face_cascade.detectMultiScale(gray, 1.3, 5)
    for (x,y,w,h)
in faces:
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,255,0),2)
    roi_gray = gray[y:y+h, x:x+w]    roi_color =
    img[y:y+h, x:x+w]

cv2_imshow(img)
```






## Face Detection in Python, OpenCV - Using a Webcam

```
import cv2
face_cascade =
    cv2.CascadeClassifier('haarcascade_frontalface_default.xml')
cap = cv2.VideoCapture(0)

while 1:
    ret, img = cap.read()
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray, 1.3, 5)

    for (x,y,w,h) in faces:
        cv2.rectangle(img,(x,y),(x+w,y+h),(255,255,0),2)
        roi_gray = gray[y:y+h, x:x+w]
        roi_color =
img[y:y+h, x:x+w]
        cv2.imshow('img',img)
        k = cv2.waitKey(30) &
0xff
        if k == 27:
            break
    cap.release()
    cv2.destroyAllWindows()
```



## Week 9 & 10

# Handling Images with Tensorflow over CoLab

## Loading Data Set (Brain Tumor Dataset)

```
import keras
from keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt import numpy as np
import cv2
import tensorflow as tf training_datagen

= ImageDataGenerator()

training_data_path = "/content/drive/MyDrive/Data Sets/Brain Tumor
Dataset/Data"

training_data = training_datagen.flow_from_directory(
training_data_path,
target_size=(224, 224), class_mode='binary'
)

Found 3772 images belonging to 2 classes.

plot_image = plt.figure(figsize=(10,10))
plot1 =
plot_image.add_subplot(2,2,1) plot2 =
```


```

plot_image.add_subplot(2,2,2) plot3 =
plot_image.add_subplot(2,2,3) plot4 =
plot_image.add_subplot(2,2,4)

    plot1.matshow(plt.imread(training_data.filepaths[20]))
plot2.matshow(plt.imread(training_data.filepaths[128]))
plot3.matshow(plt.imread(training_data.filepaths[287]))
plot4.matshow(plt.imread(training_data.filepaths[189]))

```

<matplotlib.image.AxesImage at 0x7fba2c968b90>



## CNN Model

```

model = keras.models.Sequential([
    #CNN layers
    keras.layers.Conv2D(filters=32,
        kernel_size=3,activation='relu',
        input_shape=[224, 224, 3]),
    keras.layers.MaxPooling2D(pool_size=(2,2)),
    keras.layers.Conv2D(filters=64,activation='relu',

```

```

        kernel_size=3),
        keras.layers.MaxPooling2D(pool_size=(2,2)),
        #ANN layers for classification
keras.layers.Flatten(),
        keras.layers.Dense(units=64, activation='relu'),
keras.layers.Dropout(0.1),
keras.layers.Dense(units=128, activation='relu'),
keras.layers.Dropout(0.25),
keras.layers.Dense(units=1, activation='sigmoid')
])

```

```

model.compile(optimizer = tf.optimizers.Adam(0.0001),
loss='binary_crossentropy', metrics=['accuracy']) model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_1 (Conv2D)	(None, 109, 109, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 64)	0
flatten (Flatten)	(None, 186624)	0
dense (Dense)	(None, 64)	11944000
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 128)	8320
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 1)	129
=====		

```

Total params: 11,971,841
Trainable params: 11,971,841
Non-trainable params: 0

```

```

history = model.fit(training_data, epochs=5)

```

Epoch 1/5

```

118/118 [=====] - 585s 5s/step - loss: 3.2654
- accuracy: 0.6707

```

Epoch 2/5

```

118/118 [=====] - 249s 2s/step - loss: 0.5554
- accuracy: 0.7192

```

Epoch 3/5

```

118/118 [=====] - 249s 2s/step - loss: 0.4930
- accuracy: 0.7471
Epoch 4/5
118/118 [=====] - 249s 2s/step - loss: 0.4380
- accuracy: 0.7633
Epoch 5/5
118/118 [=====] - 246s 2s/step - loss: 0.3918
- accuracy: 0.7892 training_data

```

## ResNet50

```

from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input,
decode_predictions
from tensorflow.keras.optimizers import Adam

resnet50 = ResNet50(
    include_top=True,
    weights="imagenet",
    input_tensor=None,
    input_shape=None,
    pooling=None,    classes=1000,)

Downloading data from
https://storage.googleapis.com/tensorflow/kerasapplications/resnet/resn
et50_weights_tf_dim_ordering_tf_kernels.h5 102967424/102967424
[=====] - 1s 0us/step for layer in
resnet50.layers:   layer.trainable = False

resnet_model = keras.Sequential() resnet_model.add(resnet50)

resnet_model.add(keras.layers.Dense(units=64,activation="relu"))
resnet_model.add(keras.layers.Dense(units=128,activation="relu"))
resnet_model.add(keras.layers.Dense(units=1,activation="sigmoid"))
resnet_model.summary()

Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
=====		
resnet50 (Functional)	(None, 1000)	25636712
dense_3 (Dense)	(None, 64)	64064

dense_4 (Dense)	(None, 128)	8320
dense_5 (Dense)	(None, 1)	129

=====

Total params: 25,709,225  
Trainable params: 72,513  
Non-trainable params: 25,636,712

---

```
resnet_model.compile(optimizer = Adam(learning_rate = 0.0001),
loss='binary_crossentropy', metrics=['accuracy'])
resnet_model.fit(training_data, epochs=1)
```

```
118/118 [=====] - 557s 5s/step - loss: 0.6872
- accuracy: 0.6686 <keras.callbacks.History at
```

```
0x7fba2285e250> resnet_model.fit(training_data,
epochs=1)
```

```
118/118 [=====] - 511s 4s/step - loss: 0.6631
- accuracy: 0.6461
```

```
<keras.callbacks.History at 0x7f36b6dee250>
```

## Mobile Net

```
from tensorflow.keras.applications.mobilenet import MobileNet
```

```
mobilenet = MobileNet(
input_shape=None,
include_top=True,
weights="imagenet",
input_tensor=None,
pooling=None, classes=1000,
)
```

```
for layer in mobilenet.layers: layer.trainable =
False
```

```
mobilenet_model = keras.Sequential() mobilenet_model.add(mobilenet)

mobilenet_model.add(keras.layers.Dense(units=64,activation="relu"))
mobilenet_model.add(keras.layers.Dense(units=128,activation="relu"))
mobilenet_model.add(keras.layers.Dense(units=1,activation="sigmoid"))
mobilenet_model.summary()
```



```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
=====		
mobilenet_1.00_224 (Functional)	(None, 1000)	4253864
dense_9 (Dense)	(None, 64)	64064
dense_10 (Dense)	(None, 128)	8320
dense_11 (Dense)	(None, 1)	129
=====		

```
Total params: 4,326,377
```

```
Trainable params: 4,304,489
```

```
Non-trainable params: 21,888
```

```
mobilenet_model.compile(optimizer = Adam(learning_rate = 0.0001),  
loss='binary_crossentropy', metrics= ['accuracy'])
```

```
mobilenet_model.fit(training_data, epochs=1)
```

```
118/118 [=====] - 627s 5s/step - loss: 0.6453  
- accuracy: 0.8598 <keras.callbacks.History
```

```
at 0x7f36b4cabbd0>
```

```
mobilenet_model.fit(training_data, epochs=1)
```

## VGG16

```
from tensorflow.keras.applications.vgg16 import VGG16
```

```
vgg16 = VGG16(  
input_shape=None,  
include_top=True,  
weights="imagenet",  
input_tensor=None,  
pooling=None, classes=1000,  
)
```

```
Downloading data from
```

```
https://storage.googleapis.com/tensorflow/kerasapplications/vgg16/vgg16  
_weights_tf_dim_ordering_tf_kernels.h5
```

```
553467904/553467096 [=====] - 5s 0us/step
```

```
553476096/553467096 [=====] - 5s 0us/step
```

```

for layer in vgg16.layers: layer.trainable =
False

vgg16_model = keras.Sequential() vgg16_model.add(vgg16)

vgg16_model.add(keras.layers.Dense(units=64,activation="relu"))
vgg16_model.add(keras.layers.Dense(units=128,activation="relu"))
vgg16_model.add(keras.layers.Dense(units=1,activation="sigmoid"))
vgg16_model.summary()

Model: "sequential_5"

```

---

Layer (type)	Output Shape	Param #
===== vgg16		
(Functional)	(None, 1000)	138357544
dense_12 (Dense)	(None, 64)	64064
dense_13 (Dense)	(None, 128)	8320
dense_14 (Dense)	(None, 1)	129
=====		
Total params: 138,430,057		
Trainable params: 138,430,057		
Non-trainable params: 0		

---

```

vgg16_model.compile(optimizer = Adam(learning_rate = 0.0001),
loss='binary_crossentropy', metrics=['accuracy'])
vgg16_model.fit(training_data, epochs=1)

102/118 [=====>.....] - ETA: 14:05 - loss: 0.6897 -
accuracy: 0.5494

```

## Converting normal Image to cartoon image using OpenCV & Python

```

import cv2 import numpy as np
import matplotlib.pyplot as plt

img = cv2.imread("Lena.png")
img = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.figure(figsize=(6,6)) plt.imshow(img)
plt.axis("off") plt.title("Original
Image") plt.show()

```

Original Image



```
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
gray = cv2.medianBlur(gray, 5)
plt.figure(figsize=(6,6))
plt.imshow(gray,cmap="gray") plt.axis("off")
plt.title("Grayscale Image") plt.show()
```

Grayscale Image



```
edges = cv2.adaptiveThreshold(gray, 255, cv2.ADAPTIVE_THRESH_MEAN_C,
                              cv2.THRESH_BINARY, 9, 9)
plt.figure(figsize=(6,6)) plt.imshow(edges,cmap="gray")
plt.axis("off") plt.title("Edged Image") plt.show()
```

Edged Image



```
color = cv2.bilateralFilter(img, 9, 250, 250)
cartoon = cv2.bitwise_and(color, color, mask=edges)
plt.figure(figsize=(6,16))
plt.imshow(cartoon,cmap="gray") plt.axis("off")
plt.title("Cartoon Image") plt.show()
```

Cartoon Image



## Line detection in python with OpenCV | Hough line method

```
import cv2 import numpy as np img =
cv2.imread("window.jpg") gray =
```

```
cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
edges = cv2.Canny(gray, 75, 150)
lines = cv2.HoughLinesP(edges, 1, np.pi/180, 30, maxLineGap=250)
for line in lines:
    x1, y1, x2, y2 = line[0]
    cv2.line(img, (x1, y1), (x2, y2), (0, 255, 0), 1)
cv2_imshow(edges) cv2_imshow(img)
```

